

# LING/C SC/PSYC 438/538

Lecture 22  
Sandiway Fong

# Last Time

- *Gentle introduction to probability*
- **Important notions:**
  - sample space
  - events
  - rule of counting
  - weighted events: probability
  - conditional probability:  $P(A|B)$
  - *the importance of conditional probability or expectation in language*
    - Just then, the white rabbit
    - the
    - expectation is  $p(\text{rabbit}|\text{white}) > p(\text{the}|\text{white})$  (conditional)
    - but  $p(\text{the}) > p(\text{rabbit})$  (unconditional)

# Language Models and N-grams

- **given a word sequence**
  - $w_1 w_2 w_3 \dots w_n$
- **chain rule**
  - *how to compute the probability of a sequence of words*
  - $p(w_1 w_2) = p(w_1) p(w_2|w_1)$
  - $p(w_1 w_2 w_3) = p(w_1) p(w_2|w_1) p(w_3|w_1 w_2)$
  - ...
  - $p(w_1 w_2 w_3 \dots w_n) = p(w_1) p(w_2|w_1) p(w_3|w_1 w_2) \dots p(w_n|w_1 \dots w_{n-2} w_{n-1})$
- **note**
  - It's not easy to collect (meaningful) statistics on  $p(w_n|w_{n-1} w_{n-2} \dots w_1)$  for all possible word sequences

# Language Models and N-grams

- **Given a word sequence**
  - $w_1 w_2 w_3 \dots w_n$
- **Bigram approximation**
  - *just look at the previous word only (not all the proceedings words)*
  - **Markov Assumption: finite length history**
  - 1st order Markov Model
  - $p(w_1 w_2 w_3 \dots w_n) = p(w_1) p(w_2|w_1) p(w_3|w_1 w_2) \dots p(w_n|w_1 \dots w_{n-3} w_{n-2} w_{n-1})$
  - $p(w_1 w_2 w_3 \dots w_n) \approx p(w_1) p(w_2|w_1) p(w_3|w_2) \dots p(w_n|w_{n-1})$
- **note**
  - $p(w_n|w_{n-1})$  is a lot easier to collect data for (and thus estimate well) than  $p(w_n|w_1 \dots w_{n-2} w_{n-1})$

# Language Models and N-grams

- **Trigram approximation**

- 2nd order Markov Model

- *just look at the preceding two words only*

- $p(w_1 w_2 w_3 w_4 \dots w_n) = p(w_1) p(w_2|w_1) p(w_3|w_1 w_2) p(w_4|w_1 w_2 w_3) \dots p(w_n | w_1 \dots w_{n-3} w_{n-2} w_{n-1})$

- $p(w_1 w_2 w_3 \dots w_n) \approx p(w_1) p(w_2|w_1) p(w_3|w_1 w_2) p(w_4|w_2 w_3) \dots p(w_n | w_{n-2} w_{n-1})$

- **note**

- $p(w_n | w_{n-2} w_{n-1})$  is a lot easier to estimate well than  $p(w_n | w_1 \dots w_{n-2} w_{n-1})$  but harder than  $p(w_n | w_{n-1})$

# Language Models and N-grams

- **estimating from corpora**

- *how to compute bigram probabilities*

- $p(w_n|w_{n-1}) = f(w_{n-1}w_n)/f(w_{n-1}w)$        $w$  is any word

- Since  $f(w_{n-1}w) = f(w_{n-1})$        $f(w_{n-1}) = \text{unigram frequency for } w_{n-1}$

- $p(w_n|w_{n-1}) = f(w_{n-1}w_n)/f(w_{n-1})$       **relative frequency**


- **Note:**

- The technique of estimating (true) probabilities using a relative frequency measure over a training corpus is known as **maximum likelihood estimation** (MLE)

# Language Models and N-grams

- **Typical Practice:**
  - Logprob calculations used
- **Question:**
  - *Why sum negative log of prob*
- **Answer (Part 2):**
  - $A = BC$
  - $\log(A) = \log(B) + \log(C)$
  - probabilities are in range  $(0, 1]$
  - Note:
    - want probabilities to be non-zero
    - $\log(0) = -\infty$
  - log of probabilities will be negative (up to 0)
  - take negative to make them positive

*log function*



region of  
interest

# Motivation for smoothing

- **Smoothing:** *avoid zero probability estimates*
- Consider

$$p(w_1 w_2 w_3 \dots w_n) \approx p(w_1) p(w_2|w_1) p(w_3|w_2) \dots p(w_n|w_{n-1})$$

- what happens when any individual probability component is zero?
  - **Arithmetic multiplication law:**  $0 \times X = 0$
  - ***very brittle!***
- even in a very large corpus, many possible n-grams over vocabulary space will have zero frequency
  - *particularly so for larger n-grams*



# Language Models and N-grams

- **Example:**

$W_{n-1}$

|         | I  | want | to  | eat | Chinese | food | lunch |
|---------|----|------|-----|-----|---------|------|-------|
| I       | 8  | 1087 | 0   | 13  | 0       | 0    | 0     |
| want    | 3  | 0    | 786 | 0   | 6       | 8    | 6     |
| to      | 3  | 0    | 10  | 860 | 3       | 0    | 12    |
| eat     | 0  | 0    | 2   | 0   | 19      | 2    | 52    |
| Chinese | 2  | 0    | 0   | 0   | 0       | 120  | 1     |
| food    | 19 | 0    | 17  | 0   | 0       | 0    | 0     |
| lunch   | 4  | 0    | 0   | 0   | 0       | 1    | 0     |

$W_n$

$w_{n-1}w_n$  bigram frequencies

|         |      |
|---------|------|
| I       | 3437 |
| want    | 1215 |
| to      | 3256 |
| eat     | 938  |
| Chinese | 213  |
| food    | 1506 |
| lunch   | 459  |

unigram frequencies

**Figure 6.4** Bigram counts for seven of the words (out of 1616 total word types) in the Berkeley Restaurant Project corpus of  $\approx 10,000$  sentences.

$W_{n-1}$

|         | I      | want | to    | eat   | Chinese | food  | lunch |
|---------|--------|------|-------|-------|---------|-------|-------|
| I       | .0023  | .32  | 0     | .0038 | 0       | 0     | 0     |
| want    | .0025  | 0    | .65   | 0     | .0049   | .0066 | .0049 |
| to      | .00092 | 0    | .0031 | .26   | .00092  | 0     | .0037 |
| eat     | 0      | 0    | .0021 | 0     | .020    | .0021 | .055  |
| Chinese | .0094  | 0    | 0     | 0     | 0       | .56   | .0047 |
| food    | .013   | 0    | .011  | 0     | 0       | 0     | 0     |
| lunch   | .0087  | 0    | 0     | 0     | 0       | .0022 | 0     |

$W_n$

bigram probabilities

**sparse matrix**

zeros render probabilities unusable

*(we'll need to add fudge factors - i.e. do **smoothing**)*

**Figure 6.5** Bigram probabilities for seven of the words (out of 1616 total word types) in the Berkeley Restaurant Project corpus of  $\approx 10,000$  sentences.

# Smoothing and N-grams

- **sparse dataset means zeros are a problem**
  - Zero probabilities are a problem
    - $p(w_1 w_2 w_3 \dots w_n) \approx p(w_1) p(w_2|w_1) p(w_3|w_2) \dots p(w_n|w_{n-1})$  **bigram model**
    - *one zero and the whole product is zero*
  - Zero frequencies are a problem
    - $p(w_n|w_{n-1}) = f(w_{n-1}w_n)/f(w_{n-1})$  **relative frequency**
    - *bigram  $f(w_{n-1}w_n)$  doesn't exist in dataset*
- **smoothing**
  - refers to ways of assigning zero probability n-grams a non-zero value
  - we'll look at two ways here (**just one of them today**)

# Smoothing and N-grams

- **Add-One Smoothing**

- add 1 to **all** frequency counts
- *simple and no more zeros (but there are better methods)*

- **unigram**

- $p(w) = f(w)/N$  (before Add-One)
  - $N$  = size of corpus
- $p(w) = (f(w)+1)/(N+V)$  (with Add-One)
- $f^*(w) = (f(w)+1)*N/(N+V)$  (with Add-One)
  - $V$  = number of distinct words in corpus
  - $N/(N+V)$  normalization factor adjusting for the effective increase in the corpus size caused by Add-One

must rescale so  
that total  
probability  
mass stays at 1

- **bigram**

- $p(w_n|w_{n-1}) = f(w_{n-1}w_n)/f(w_{n-1})$  (before Add-One)
- $p(w_n|w_{n-1}) = (f(w_{n-1}w_n)+1)/(f(w_{n-1})+V)$  (after Add-One)
- $f^*(w_{n-1}w_n) = (f(w_{n-1}w_n)+1)*f(w_{n-1})/(f(w_{n-1})+V)$  (after Add-One)

# Smoothing and N-grams

- **Add-One Smoothing**
  - add 1 to all frequency counts
- **bigram**
  - $p(w_n|w_{n-1}) = (f(w_{n-1}w_n)+1)/(f(w_{n-1})+V)$
  - $(f(w_{n-1}w_n)+1) * f(w_{n-1}) / (f(w_{n-1})+V)$
- **frequencies**

|         | I  | want | to  | eat | Chinese | food | lunch |
|---------|----|------|-----|-----|---------|------|-------|
| I       | 8  | 1087 | 0   | 13  | 0       | 0    | 0     |
| want    | 3  | 0    | 786 | 0   | 6       | 8    | 6     |
| to      | 3  | 0    | 10  | 860 | 3       | 0    | 12    |
| eat     | 0  | 0    | 2   | 0   | 19      | 2    | 52    |
| Chinese | 2  | 0    | 0   | 0   | 0       | 120  | 1     |
| food    | 19 | 0    | 17  | 0   | 0       | 0    | 0     |
| lunch   | 4  | 0    | 0   | 0   | 0       | 1    | 0     |

= figure 6.4

|         | I    | want   | to     | eat    | Chinese | food  | lunch |
|---------|------|--------|--------|--------|---------|-------|-------|
| I       | 6.12 | 740.05 | 0.68   | 9.52   | 0.68    | 0.68  | 0.68  |
| want    | 1.72 | 0.43   | 337.76 | 0.43   | 3.00    | 3.86  | 3.00  |
| to      | 2.67 | 0.67   | 7.35   | 575.41 | 2.67    | 0.67  | 8.69  |
| eat     | 0.37 | 0.37   | 1.10   | 0.37   | 7.35    | 1.10  | 19.47 |
| Chinese | 0.35 | 0.12   | 0.12   | 0.12   | 0.12    | 14.09 | 0.23  |
| food    | 9.65 | 0.48   | 8.68   | 0.48   | 0.48    | 0.48  | 0.48  |
| lunch   | 1.11 | 0.22   | 0.22   | 0.22   | 0.22    | 0.44  | 0.22  |

= figure 6.8

## Remarks:

*perturbation problem*

add-one causes large changes in some frequencies due to relative size of  $V$  (1616)

*want to: 786  $\Rightarrow$  338*

# Smoothing and N-grams

- **Add-One Smoothing**
  - *add 1 to all frequency counts*
- **bigram**
  - $p(w_n|w_{n-1}) = (f(w_{n-1}w_n)+1)/(f(w_{n-1})+V)$
  - $(f(w_{n-1}w_n)+1) * f(w_{n-1}) / (f(w_{n-1})+V)$
- **Probabilities**

**Remarks:**  
*perturbation problem*

similar changes in  
probabilities

|         | I       | want    | to      | eat     | Chinese | food    | lunch   |
|---------|---------|---------|---------|---------|---------|---------|---------|
| I       | 0.00233 | 0.31626 | 0.00000 | 0.00378 | 0.00000 | 0.00000 | 0.00000 |
| want    | 0.00247 | 0.00000 | 0.64691 | 0.00000 | 0.00494 | 0.00658 | 0.00494 |
| to      | 0.00092 | 0.00000 | 0.00307 | 0.26413 | 0.00092 | 0.00000 | 0.00369 |
| eat     | 0.00000 | 0.00000 | 0.00213 | 0.00000 | 0.02026 | 0.00213 | 0.05544 |
| Chinese | 0.00939 | 0.00000 | 0.00000 | 0.00000 | 0.00000 | 0.56338 | 0.00469 |
| food    | 0.01262 | 0.00000 | 0.01129 | 0.00000 | 0.00000 | 0.00000 | 0.00000 |
| lunch   | 0.00871 | 0.00000 | 0.00000 | 0.00000 | 0.00000 | 0.00218 | 0.00000 |


= figure 6.5

|         | I       | want    | to      | eat     | Chinese | food    | lunch   |
|---------|---------|---------|---------|---------|---------|---------|---------|
| I       | 0.00178 | 0.21532 | 0.00020 | 0.00277 | 0.00020 | 0.00020 | 0.00020 |
| want    | 0.00141 | 0.00035 | 0.27799 | 0.00035 | 0.00247 | 0.00318 | 0.00247 |
| to      | 0.00082 | 0.00021 | 0.00226 | 0.17672 | 0.00082 | 0.00021 | 0.00267 |
| eat     | 0.00039 | 0.00039 | 0.00117 | 0.00039 | 0.00783 | 0.00117 | 0.02075 |
| Chinese | 0.00164 | 0.00055 | 0.00055 | 0.00055 | 0.00055 | 0.06616 | 0.00109 |
| food    | 0.00641 | 0.00032 | 0.00577 | 0.00032 | 0.00032 | 0.00032 | 0.00032 |
| lunch   | 0.00241 | 0.00048 | 0.00048 | 0.00048 | 0.00048 | 0.00096 | 0.00048 |

= figure 6.7


# Smoothing and N-grams

- *let's illustrate the problem*
  - take the bigram case:
  - $w_{n-1}w_n$
  - $p(w_n|w_{n-1}) = f(w_{n-1}w_n)/f(w_{n-1})$
  - suppose there are cases
  - $w_{n-1}w_1^0$  that don't occur in the corpus




# Smoothing and N-grams

- *add-one*
  - “give everyone 1”



# Smoothing and N-grams

- *add-one*
  - “give everyone 1”
- *redistribution of probability mass*
  - $p(w_n|w_{n-1}) = (f(w_{n-1}w_n)+1)/(f(w_{n-1})+V)$





# Smoothing and N-grams

- Excel spreadsheet available
  - addone.xls

# Smoothing and N-grams

- **Witten-Bell Smoothing**

- equate zero frequency items with frequency 1 items
- use frequency of things seen once to estimate frequency of things we haven't seen yet
- *smaller impact than Add-One*

- **unigram**

- a zero frequency word (unigram) is “an event that hasn't happened yet”
- count the number of (*different*) words (T) we've observed in the corpus
- $p(w) = T/(Z*(N+T))$ 
  - w is a word with zero frequency
  - Z = number of zero frequency words
  - N = size of corpus

- **bigram**

- $p(w_n|w_{n-1}) = f(w_{n-1}w_n)/f(w_{n-1})$  (*original*)
- $p(w_n|w_{n-1}) = T(w_{n-1})/(Z(w_{n-1})*(T(w_{n-1})+N(w_{n-1})))$  for zero bigrams (*after* Witten-Bell)
  - $T(w_{n-1})$  = number of (seen) bigrams beginning with  $w_{n-1}$
  - $Z(w_{n-1})$  = number of unseen bigrams beginning with  $w_{n-1}$
  - $Z(w_{n-1}) = (possible\ bigrams\ beginning\ with\ w_{n-1}) - (the\ ones\ we've\ seen)$
  - $Z(w_{n-1}) = V - T(w_{n-1})$
- $T(w_{n-1})/Z(w_{n-1}) * f(w_{n-1})/(f(w_{n-1})+T(w_{n-1}))$  estimated zero bigram frequency
- $p(w_n|w_{n-1}) = f(w_{n-1}w_n)/(f(w_{n-1})+T(w_{n-1}))$  for non-zero bigrams (*after* Witten-Bell)

# Smoothing and N-grams

- **Witten-Bell Smoothing**

- use frequency of things seen once to estimate frequency of things we haven't seen yet

- **bigram**

- $T(w_{n-1}) / Z(w_{n-1}) * f(w_{n-1}) / (f(w_{n-1}) + T(w_{n-1}))$  estimated zero bigram frequency
  - $T(w_{n-1})$  = number of bigrams beginning with  $w_{n-1}$
  - $Z(w_{n-1})$  = number of unseen bigrams beginning with  $w_{n-1}$

|         | I  | want | to  | eat | Chinese | food | lunch |
|---------|----|------|-----|-----|---------|------|-------|
| I       | 8  | 1087 | 0   | 13  | 0       | 0    | 0     |
| want    | 3  | 0    | 786 | 0   | 6       | 8    | 6     |
| to      | 3  | 0    | 10  | 860 | 3       | 0    | 12    |
| eat     | 0  | 0    | 2   | 0   | 19      | 2    | 52    |
| Chinese | 2  | 0    | 0   | 0   | 0       | 120  | 1     |
| food    | 19 | 0    | 17  | 0   | 0       | 0    | 0     |
| lunch   | 4  | 0    | 0   | 0   | 0       | 1    | 0     |

= figure 6.4

**Remark:**  
*smaller changes*

|         | I      | want     | to      | eat     | Chinese | food    | lunch  |
|---------|--------|----------|---------|---------|---------|---------|--------|
| I       | 7.785  | 1057.763 | 0.061   | 12.650  | 0.061   | 0.061   | 0.061  |
| want    | 2.823  | 0.046    | 739.729 | 0.046   | 5.647   | 7.529   | 5.647  |
| to      | 2.885  | 0.084    | 9.616   | 826.982 | 2.885   | 0.084   | 11.539 |
| eat     | 0.073  | 0.073    | 1.766   | 0.073   | 16.782  | 1.766   | 45.928 |
| Chinese | 1.828  | 0.011    | 0.011   | 0.011   | 0.011   | 109.700 | 0.914  |
| food    | 18.019 | 0.051    | 16.122  | 0.051   | 0.051   | 0.051   | 0.051  |
| lunch   | 3.643  | 0.026    | 0.026   | 0.026   | 0.026   | 0.911   | 0.026  |

= figure 6.9

# Smoothing and N-grams

- Witten-Bell excel spreadsheet:
  - `wb.xls`

# Smoothing and N-grams

- Witten-Bell Smoothing
- *Implementation (Excel)*
  - one line formula

sheet1

|         | I  | want | to  | eat | Chinese | food | lunch |      |
|---------|----|------|-----|-----|---------|------|-------|------|
| I       | 8  | 1087 | 0   | 13  | 0       | 0    | 0     | 3437 |
| want    | 3  | 0    | 786 | 0   | 6       | 8    | 6     | 1215 |
| to      | 3  | 0    | 10  | 860 | 3       | 0    | 12    | 3256 |
| eat     | 0  | 0    | 2   | 0   | 19      | 2    | 52    | 938  |
| Chinese | 2  | 0    | 0   | 0   | 0       | 120  | 1     | 213  |
| food    | 19 | 0    | 17  | 0   | 0       | 0    | 0     | 1506 |
| lunch   | 4  | 0    | 0   | 0   | 0       | 1    | 0     | 459  |

B2

=IF(Sheet1!B2=0,\$J2/\$K2\*(Sheet1!\$J2/(Sheet1!\$J2+\$J2)),Sheet1!B2\*(Sheet1!\$J2/(Sheet1!\$J2+\$J2)))

Workbook2

sheet2: rescaled sheet1

|         | I      | want     | to      | eat     | Chinese | food    | lunch  |  |
|---------|--------|----------|---------|---------|---------|---------|--------|--|
| I       | 7.785  | 1057.763 | 0.061   | 12.650  | 0.061   | 0.061   | 0.061  |  |
| want    | 2.823  | 0.046    | 739.729 | 0.046   | 5.647   | 7.529   | 5.647  |  |
| to      | 2.885  | 0.084    | 9.616   | 826.982 | 2.885   | 0.084   | 11.539 |  |
| eat     | 0.073  | 0.073    | 1.766   | 0.073   | 16.782  | 1.766   | 45.928 |  |
| Chinese | 1.828  | 0.011    | 0.011   | 0.011   | 0.011   | 109.700 | 0.914  |  |
| food    | 18.019 | 0.051    | 16.122  | 0.051   | 0.051   | 0.051   | 0.051  |  |
| lunch   | 3.643  | 0.026    | 0.026   | 0.026   | 0.026   | 0.911   | 0.026  |  |

| T   | Z    |
|-----|------|
| 95  | 1521 |
| 76  | 1540 |
| 130 | 1486 |
| 124 | 1492 |
| 20  | 1596 |
| 82  | 1534 |
| 45  | 1571 |

# Language Models and N-grams

- **N-gram models**

- they're technically easy to compute
  - *(in the sense that lots of training data are available)*
- but just how good are these n-gram language models?
- and what can they show us about language?

# Language Models and N-grams

## approximating Shakespeare

- generate random sentences using n-grams
- train on complete *Works of Shakespeare*
- **Unigram** (pick random, unconnected words)
  - (a) To him swallowed confess hear both. Which. Of save on trail for are ay device and rote life have
  - (b) Every enter now severally so, let
  - (c) Hill he late speaks; or! a more to leg less first you enter
  - (d) Will rash been and by I the me loves gentle me not slavish page, the and hour; ill let
  - (e) Are where exeunt and sighs have rise excellency took of.. Sleep knave we. near; vile like
- **Bigram**
  - (a) What means, sir. I confess she? then all sorts, he is trim, captain.
  - (b) Why dost stand forth thy canopy, forsooth; he is this palpable hit the King Henry. Live king. Follow.
  - (c) What we, hath got so she that I rest and sent to scold and nature bankrupt, nor the first gentleman?
  - (d) Enter Menenius, if it so many good direction found'st thou art a strong upon command of fear not a liberal largess given away, Falstaff! Exeunt
  - (e) Thou whoreson chops. Consumption catch your dearest friend, well, and I know where many mouths upon my undoing all but be, how soon, then; we'll execute upon my love's bonds and we do you will?
  - (f) The world shall- my lord!

# Language Models and N-grams

- Approximating Shakespeare (section 6.2)
  - generate random sentences using n-grams
  - train on complete *Works of Shakespeare*
- **Trigram**
  - (a) Sweet prince, Falstaff shall die. Harry of Monmouth's grave.
  - (b) This shall forbid it should be branded, if renown made it empty.
  - (c) What is't that cried?
  - (d) Indeed the duke; and had a very good friend.
  - (e) Fly, and will rid me these news of price. Therefore the sadness of parting, as they say, 'tis done.
  - (f) The sweet! How many then shall posthumus end his miseries.
- **Quadrigram**
  - (a) King Henry. What! I will go seek the traitor Gloucester. Exeunt some of the watch. A great banquet serv'd in;
  - (b) Will you not tell me who I am?
  - (c) It cannot be but so.
  - (d) Indeed the short and the long. Marry, 'tis a noble Lepidus.
  - (e) They say all lovers swear more performance than they are wont to keep obliged faith unforfeited!
  - (f) Enter Leonato's brother Antonio, and the rest, but seek the weary beds of people sick.

**Remarks:**  
*dataset size problem*

training set is small  
884,647 words  
29,066 different words

$29,066^2 = 844,832,356$   
possible bigrams

for the random sentence generator, this means very limited choices for possible continuations, which means program can't be very innovative for higher  $n$



# Possible Application?

# Language Models and N-grams

- **Aside:** <http://hemispheresmagazine.com/contests/2004/intro.htm>

**HEMISPHERES** | [Search](#) | [Reader Service](#) | [Media Kit](#) | [Contact Us](#) | [Special Ad Sections](#) | [Contests](#)

**CONTENTS**  
[How to enter the contest](#)

**2004**  
[Faulkner parodies](#)  
[Hemingway parodies](#)

**2003**  
[Faulkner parodies](#)  
[Hemingway parodies](#)

**2002**  
[Faulkner parodies](#)  
[Hemingway parodies](#)

**2001**  
[Faulkner parodies](#)  
[Hemingway parodies](#)

**2000**  
[Faulkner parodies](#)  
[Hemingway parodies](#)

**CONTESTS**

**The 2004 Faux Faulkner and Imitation Hemingway Contests**  
[Faulkner parodies](#) • [Hemingway parodies](#)

**As I Lay Laughing: Winners of the Faux Faulkner and Imitation Hemingway Parody Contests**

For the fifth consecutive year, hemispheres is delighted to publish the winning entries from the annual Faux Faulkner and Imitation Hemingway parody contests. Hundreds of entries poured in from across the globe, and from them our judges plucked a Faulknerian tribute to The Three Stooges and a Hemingwayesque dating-service personal ad. Judges this year included Arthur Schlesinger Jr., John Berendt, Roy Blount Jr., Digby Diehl, Barry Hannah, and Jill McCorkle. Sadly, George Plimpton—an enthusiastic Faulkner judge for all 14 years of the competition—died prior to this year’s contest. In the 15th annual Faulkner contest, writer David Sheffield was the winner with his “As I Lay Kvetching.” And author Deborah Ein’s “Dear Scott Letter” took the 24th annual Hemingway prize. And now, let the laughing begin.

# Language Models and N-grams

- **N-gram models + smoothing**
  - one consequence of smoothing is that
  - every possible concatenation or sequence of words has a non-zero probability

# Colorless green ideas

- **examples**
  - (1) **colorless green ideas** sleep furiously
  - (2) furiously sleep ideas green colorless
- **Chomsky (1957):**
  - . . . It is fair to assume that neither sentence (1) nor (2) (nor indeed any part of these sentences) has ever occurred in an English discourse. Hence, **in any statistical model for grammaticality**, these sentences will be ruled out on identical grounds as equally 'remote' from English. Yet (1), though nonsensical, is grammatical, while (2) is not.
- **idea**
  - (1) is syntactically valid, (2) is word salad
- Statistical Experiment (Pereira 2002)

# Colorless green ideas

- **examples**
  - (1) **colorless green ideas** sleep furiously
  - (2) furiously sleep ideas green colorless
- Statistical Experiment (Pereira 2002)

|           |       |
|-----------|-------|
| $w_{i-1}$ | $w_i$ |
|-----------|-------|

$$p(w_1 \cdots w_n) = p(w_1) \prod_{i=2}^n p(w_i | w_{i-1}) \quad . \quad \text{bigram language model}$$

Using this estimate for the probability of a string and an aggregate model with  $C = 16$  trained on newspaper text using the expectation-maximization (EM) method (Dempster, Laird, & Rubin, 1977), we find that

$$\frac{p(\text{Colorless green ideas sleep furiously})}{p(\text{Furiously sleep ideas green colorless})} \approx 2 \times 10^5 \quad .$$

Thus, a suitably constrained statistical model, even a very simple one, can meet Chomsky's particular challenge.

# Interesting things to Google™

- **example**
  - **colorless green ideas** sleep furiously
- **First hit**

## Web

### **Colorless green ideas sleep furiously**

Chomsky's famous sentence '**Colorless green ideas** sleep furiously' is examined and is shown to be a specimen of irony rather being meaningless.

[home.tiac.net/~cri/1997/chomsky.html](http://home.tiac.net/~cri/1997/chomsky.html) - 4k - [Cached](#) - [Similar pages](#)

# Interesting things to Google™

- **example**
  - **colorless green ideas** sleep furiously
- **first hit**
  - **compositional semantics**
  - a **green idea** is, according to well established usage of the word "green" is one that is an idea that is **new and untried**.
  - again, a **colorless idea** is one **without vividness, dull and unexciting**.
  - so it follows that a **colorless green idea** is a **new, untried idea that is without vividness, dull and unexciting**.
  - to **sleep** is, among other things, is to be in a state of dormancy or inactivity, or in a state of unconsciousness.
  - to **sleep furiously** may seem a puzzling turn of phrase but one reflects that **the mind in sleep often indeed moves furiously with ideas and images flickering in and out**.

# Interesting things to Google™

- **example**
  - **colorless green ideas** sleep furiously
- **another hit:** (*a story*)
  - "So this is our ranking system," said Chomsky. "As you can see, the highest rank is yellow."
  - "And the **new ideas**?"
  - "The **green ones**? Oh, the **green ones don't get a color until they've had some seasoning**. These ones, anyway, are still too angry. **Even when they're asleep, they're furious**. We've had to kick them out of the dormitories - they're just unmanageable."
  - "So where are they?"
  - "Look," said Chomsky, and pointed out of the window. There below, on the lawn, the colorless green ideas slept, furiously.



# More on N-grams

- How to degrade gracefully when we don't have evidence
  - Backoff
  - Deleted Interpolation
- N-grams and Spelling Correction

# Backoff

- **idea**
  - Hierarchy of approximations
  - trigram > bigram > unigram
  - *degrade gracefully*
- Given a word sequence fragment:
  - $\dots w_{n-2} w_{n-1} w_n \dots$
- **preference rule**
  1.  $p(w_n | w_{n-2} w_{n-1})$  if  $f(w_{n-2} w_{n-1} w_n) \neq 0$
  2.  $\alpha_1 p(w_n | w_{n-1})$  if  $f(w_{n-1} w_n) \neq 0$
  3.  $\alpha_2 p(w_n)$
- notes:
  - $\alpha_1$  and  $\alpha_2$  are fudge factors to ensure that probabilities still sum to 1

# Backoff

- **preference rule**

1.  $p(w_n | w_{n-2} w_{n-1})$  if  $f(w_{n-2} w_{n-1} w_n) \neq 0$
2.  $\alpha_1 p(w_n | w_{n-1})$  if  $f(w_{n-1} w_n) \neq 0$
3.  $\alpha_2 p(w_n)$

- **problem**


- if  $f(w_{n-2} w_{n-1} w_n) = 0$ , we use one of the estimates from (2) or (3)
- assume the backoff value is non-zero
- then we are introducing non-zero probability for  $p(w_n | w_{n-2} w_{n-1})$  – which is zero in the corpus
- then this adds “probability mass” to  $p(w_n | w_{n-2} w_{n-1})$  which is not in the original system
- *therefore, we have to be careful to juggle the probabilities to still sum to 1*

# Deleted Interpolation

– *fundamental idea of interpolation*

– **equation:** (*trigram*)

- $p(w_n | w_{n-2} w_{n-1}) =$ 
  - $\lambda_1 p(w_n | w_{n-2} w_{n-1}) +$
  - $\lambda_2 p(w_n | w_{n-2}) +$
  - $\lambda_3 p(w_n)$



– **Note:**

- $\lambda_1$ ,  $\lambda_2$  and  $\lambda_3$  are fudge factors to ensure that probabilities still sum to 1